Artificial Neural Networks & Deep Learning Project - Image Segmentation

Team 10520637_10527649

December 17, 2019

1   Lab Model

We implemented a first model based on the notebook provided during the lab session.

Setup summary:

- Validation Split done outside of ImageDataGenerator

- Batch Size = 16

- Encoder depth = 4

- Encoder block: Conv2D+ReLU → Conv2D+ReLU → MaxPool2D

- Last Encoder layer: Conv2D+ReLU → Conv2D+ReLU → MaxPool2D

- Decoder depth = 4

- Decoder block: UpSampling2D → Conv2D+ReLU → Conv2D+ReLU

- Loss function: SparseCategoricalCrossEntropy

- Optimizer: Adam(learning_rate = `1e-3`)

This turned out in `train_loss=0.2949`, `valid_loss=0.3012` and `kaggle_score=0.21515`.


Figure 1: CNN `loss` plot

Figure 2: CNN `My_IoU` plot

The graphs show a sudden decrease because the training was interrupted before terminating.

1.1   First Model with concatenation

We tried to add concatenation to the model provided during the lab session. The setup did not change noticeably. This turned out in `train_loss=0.250`, `valid_loss=0.300` and `kaggle_score=0.37682`.


Figure 3: CNN with concatenation `loss` plot

Figure 4: CNN with concatenation`My_IoU` plot

2   U-Net

We realized we were not applying concatenation correctly so we reimplemented a U-Net style model using Keras Functional API, explicitly declaring each layer. We used the same setup as before. We realized the model was not performing as good as we expected because the graphs showed that the model was not learning. This started a series of attempts to change the model hyperparameters. Our best attempt resulted in `train_loss=0.21`, `val_loss=0.31` and `kaggle_score=0.39594`


Figure 5: U-Net `loss` plot

Figure 6: U-Net `My_IoU` plot

3   Failing attempts

The main issue with the model was that the training went straight to a local optimum with `train_loss=0` and `train_my_IoU=0`.
We noticed that the `num_classes` parameter was used to define the number of filters in the last layer and since we needed just one image as prediction output we changed to `num_classes=1`.
Adding Stochastic Gradient Descent to model optimizers slightly improved the model performances since it didn't immediately go down to `my_IoU=0`.

3.1 Loss Functions

We thought the model was going into a local optimum due to the loss function we were using. To overcome this problem we tried different loss functions.

3.1.1 Binary Cross Entropy

The Cross Entropy loss function measures the accuracy of the network over the single pixel.

$$\text{CE}(p, \hat{p}) = -\left(p \log(\hat{p}) + (1 - p) \log(1 - \hat{p})\right)$$

3.1.2 Dice Loss

The Dice loss gives a feedback closely related to IoU (Intersection over Union). It measures the average prediction accuracy over the target instead of looking at the precision of single pixels.

$$\text{DC} = \frac{2TP}{2TP + FP + FN} = \frac{2|X \cap Y|}{|X| + |Y|}$$

- TP = True Positives

- FP = False Positives

- FN = False Negatives

- X = Prediction

- Y = Target

3.1.3 Combined Loss

We decided to use Binary Cross Entropy in combination with Dice Loss, giving the same weight to both, in order to find a model able to identify high-level features and without losing accuracy at a low-level scale.

3.2 Dataset Split

Following the advise of Tutor F. Lattari we changed our dataset split method and rolled back to the `ImageDataGenerator` built-in `validation_split` feature. This showed immediate improvements, meaning we had some errors in the code used to split the dataset.

4   Transfer Learning

All the previous attempts to improve the last model turned out in very low Kaggle scores. We inferred the model was underfitting so we tried to implement Transfer Learning.
Changes in the setup:

- Batch size = 8

- Loss Function = `combined_loss`

- Dropout layers in Decoder

4.1   VGG16

We used VGG16 as Encoder with fine-tuning of the last six layers and the U-Net Keras-Functional-API implementation Decoder mentioned before.
This model finished the training with `train_loss=0.4653`, `valid_loss=0.5639` and `kaggle_score=0.47406`.


Figure 7: VGG16 `loss` plot

Figure 8: VGG16 `My_IoU` plot

4.2   VGG16 with concatenation

We implemented VGG16 using Keras Functional API and initiliazed the model with weights imported from ImageNet. This model finished with:
`train_loss=0.4535`, `valid_loss=0.5470` and `kaggle_score=0.52884`.

Figure 9: VGG16 concatenation `loss` plot

Figure 10: VGG16 concatenation `My_IoU` plot

4.3 VGG19 with concatenation

We repeated the attempt with a more recent architecture.

This model finished with `train_loss=0.`, `valid_loss=0.` and `kaggle_score=0.`.

4.4 Xception with concatenation

At last we tried to implement the Encoder using Xception network. We launched two trainings: one with fine tuning after 116 layers, the other after 126 layers. We obtained the best results with . The training finished with `train_loss=0.`, `valid_loss=0.` and `kaggle_score=0.`.

The model with which we obtained the highest score is: